

第3-3章 点估计 Bayes估计

《统计推断》 第7章

感谢清华大学自动化系江瑞教授提供PPT

内容

- Bayes估计方法
- 共轭先验
- Bayes估计应用

How to update the degree of belief?

$$p(\theta \mid \mathbf{x}) = \frac{p(\theta)p(\mathbf{x} \mid \theta)}{p(\mathbf{x})}$$
$$= \frac{\text{prior} \times \text{likelihood}}{\text{marginal likelihood}}$$

$p(\theta)$: prior

$p(\mathbf{x} \mid \theta)$: likelihood

$p(\mathbf{x})$: marginal likelihood (evidence)

$p(\theta \mid \mathbf{x})$: posterior

Bernoulli likelihood

Let X_1, \dots, X_n be a sample from a Bernoulli (θ) population, we have the likelihood function as

$$\begin{aligned} p(\mathbf{x} \mid \theta) &= \prod_{i=1}^n \theta^{x_i} (1 - \theta)^{1-x_i} \\ &= \theta^{n_1} (1 - \theta)^{n-n_1} \end{aligned}$$

$$n_1 = \sum_{i=1}^n x_i$$

Beta prior

The likelihood is

$$p(\mathbf{x} \mid \theta) = \theta^{n_1} (1 - \theta)^{n - n_1}$$

A natural selection of the prior is the Beta distribution

$$p(\theta \mid \alpha, \beta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1} (1 - \theta)^{\beta-1}$$

The parameters in the prior are called **hyper-parameters**

Integrate out the parameter

$$\begin{aligned} p(\mathbf{x}) &= \int p(\theta) p(\mathbf{x} \mid \theta) d\theta \\ &= \int \left[\frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1} (1 - \theta)^{\beta-1} \right] \left[\theta^{n_1} (1 - \theta)^{n-n_1} \right] d\theta \\ &= \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \underbrace{\int \theta^{\alpha+n_1-1} (1 - \theta)^{\beta+n-n_1-1} d\theta}_{\substack{= \left[\frac{\Gamma(\alpha+\beta+n)}{\Gamma(\alpha+n_1)\Gamma(\beta+n-n_1)} \right]^{-1}}} \\ &= \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \frac{\Gamma(\alpha + n_1)\Gamma(\beta + n - n_1)}{\Gamma(\alpha + \beta + n)} \end{aligned}$$

Obtain the posterior

$$\begin{aligned} p(\theta \mid \mathbf{x}) &= \frac{p(\theta)p(\mathbf{x} \mid \theta)}{p(\mathbf{x})} \\ &= \frac{\left[\frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1} (1 - \theta)^{\beta-1} \right] \left[\theta^{n_1} (1 - \theta)^{n-n_1} \right]}{\frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \frac{\Gamma(\alpha + n_1)\Gamma(\beta + n - n_1)}{\Gamma(\alpha + \beta + n)}} \\ &= \frac{\Gamma(\alpha + \beta + n)}{\Gamma(\alpha + n_1)\Gamma(\beta + n - n_1)} \theta^{\alpha+n_1-1} (1 - \theta)^{\beta+n-n_1-1} \end{aligned}$$

That is to say

$$\theta \sim \text{Beta}(\alpha, \beta)$$

$$\theta \mid \mathbf{x} \sim \text{Beta}(\alpha + n_1, \beta + n - n_1)$$

Prior and posterior comes from the same parametric family of distributions, just the parameters are different (updated with the statistic from the sample)

Conjugate prior

Let \mathcal{F} denote the class of pdfs or pmfs $f(x | \theta)$ (indexed by θ). A class Π of prior distribution is a **conjugate family** of \mathcal{F} if the posterior distribution is in the class Π for all $f \in \mathcal{F}$, all priors in Π , and all $x \in \mathcal{X}$.

Derivation of a posterior distribution

- Write the likelihood function
- Select a prior distribution, generally select a conjugate prior
- Integrate out the parameter to obtain the marginal likelihood
- Derive the posterior (with the conjugate prior, the posterior has the same form as prior)

Start from a full probabilistic model

A **joint probability distribution** for all observable and unobservable quantities in a problem. The model should be consistent with knowledge about the underlying scientific problem and the data collection process

$$p(\theta \mid \mathbf{x}) = \frac{p(\theta, \mathbf{x})}{p(\mathbf{x})} \propto p(\theta, \mathbf{x}) = p(\theta)p(\mathbf{x} \mid \theta)$$

$$p(\theta, \mathbf{x}) = p(\theta)p(\mathbf{x} \mid \theta) \quad (\text{Chain rule})$$

A simple way

$$p(\theta \mid \mathbf{x}) = \frac{p(\theta)p(\mathbf{x} \mid \theta)}{p(\mathbf{x})} \propto p(\theta)p(\mathbf{x} \mid \theta)$$

$$p(\mathbf{x} \mid \theta) \propto \theta^{n_1} (1 - \theta)^{n - n_1}$$

$$p(\theta) \propto \theta^{\alpha-1} (1 - \theta)^{\beta-1}$$

$$p(\theta \mid \mathbf{x}) \propto \theta^{\alpha+n_1-1} (1 - \theta)^{\beta+n-n_1-1}$$

$$p(\theta \mid \mathbf{x}) = \kappa \theta^{\alpha+n_1-1} (1 - \theta)^{\beta+n-n_1-1}$$

$$\int p(\theta \mid \mathbf{x}) d\theta = 1 \Rightarrow \kappa = \left[\int \theta^{\alpha+n_1-1} (1 - \theta)^{\beta+n-n_1-1} d\theta \right]^{-1}$$

$$= \frac{\Gamma(\alpha + \beta + n)}{\Gamma(\alpha + n_1) \Gamma(\beta + n - n_1)}$$

Maximum *a posteriori* (MAP) estimate

Because the posterior distribution $p(\theta | \mathbf{x})$ describes the distribution of the parameter. It is reasonable to choose the parameter θ^{**} that can maximize the posterior pdf or pmf $p(\theta | \mathbf{x})$ as our best guess of θ , given the observed data. In other words,

$$\theta^{**} = \arg \max_{\theta \in \Theta} p(\theta | \mathbf{x}).$$

Equivalently,

$$\theta^{**} = \arg \max_{\theta \in \Theta} p(\theta)p(\mathbf{x} | \theta).$$

Or,

$$\theta^{**} = \arg \max_{\theta \in \Theta} \log p(\theta)p(\mathbf{x} | \theta).$$

Obviously,

$$p(\theta^{**} | \mathbf{x}) \geq p(\theta | \mathbf{x}), \text{ for any } \theta \in \Theta.$$

θ^{**} is called the *maximum a posteriori* (MAP) *estimate* of θ .

Bernoulli MAP

The posterior distribution is

$$\theta \mid \mathbf{x} \sim \text{Beta}(\alpha + n_1, \beta + n - n_1)$$

However, it is known that the point with maximum density in a $\text{Beta}(\alpha, \beta)$ distribution is

$$\frac{\alpha - 1}{\alpha + \beta - 2}$$

Therefore, the MAP estimate for the Bernoulli parameter is

$$\theta^{**} = \frac{\alpha + n_1 - 1}{\alpha + \beta + n - 2}$$

Bayes estimate

Bayes estimate

Because

$$p(\theta \mid \mathbf{x}) = \frac{p(\theta)p(\mathbf{x} \mid \theta)}{p(\mathbf{x})}$$

The **posterior mean** is a natural point estimate of θ .

Precisely, the objective of Bayesian inference is not to obtain a single value of a parameter, but the distribution of the parameter.

Bernoulli Bayes estimate

The mean of a $\text{Beta}(\alpha, \beta)$ distribution is

$$\frac{\alpha}{\alpha + \beta}$$

Therefore, the Bayes estimate for the Bernoulli parameter is

$$\hat{\theta} = \frac{\alpha + n_1}{\alpha + \beta + n}$$

Pseudo counts

$$\begin{aligned}\hat{\theta}_{\text{Bayes}} &= \frac{\alpha + n_1}{\alpha + \beta + n} \\ &= \underbrace{\left(\frac{\alpha + \beta}{\alpha + \beta + n} \right)}_{\text{Prior weight}} \underbrace{\left(\frac{\alpha}{\alpha + \beta} \right)}_{\text{Prior mean}} + \underbrace{\left(\frac{n}{\alpha + \beta + n} \right)}_{\text{Sample weight}} \underbrace{\left(\frac{n_1}{n} \right)}_{\text{Sample mean}}\end{aligned}$$

Posterior mean is the weighted average of the prior mean and the sample mean

They are different

Maximum likelihood estimate (MLE)

$$\hat{\theta}_{\text{MLE}} = \frac{n_1}{n}$$



Maximum a posteriori estimate (MAP)

$$\hat{\theta}_{\text{MAP}} = \frac{\alpha + n_1 - 1}{\alpha + \beta + n - 2}$$

Bayes estimate (mean of the posterior distribution)

$$\hat{\theta}_{\text{Bayes}} = \frac{\alpha + n_1}{\alpha + \beta + n}$$

Non-informative prior ($\alpha=1, \beta=1$)



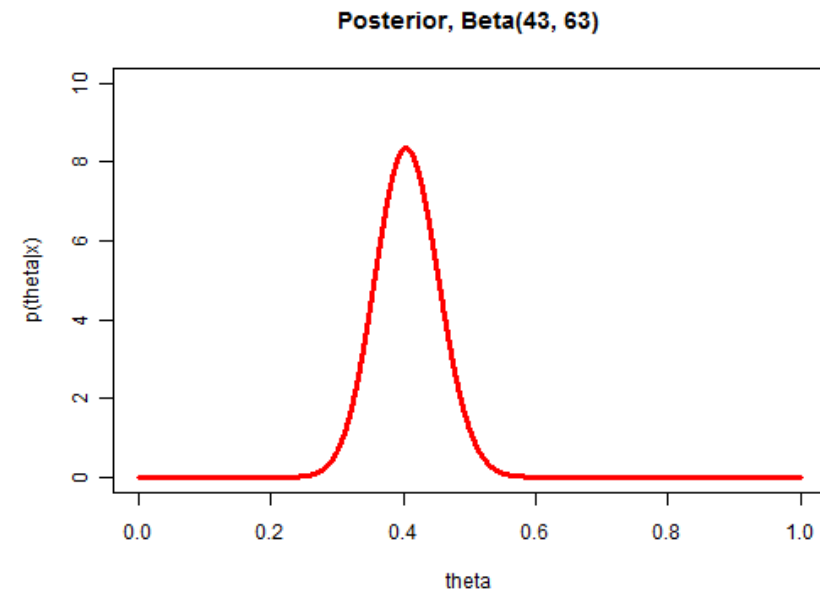

$$\hat{\theta}_{\text{MLE}} = \frac{40}{100}$$

$$\hat{\theta}_{\text{MAP}} = \frac{40}{100}$$

$$\hat{\theta}_{\text{Bayes}} = \frac{41}{101}$$

$$n = 100, n_1 = 40$$

Weak prior ($\alpha=3, \beta=3$)



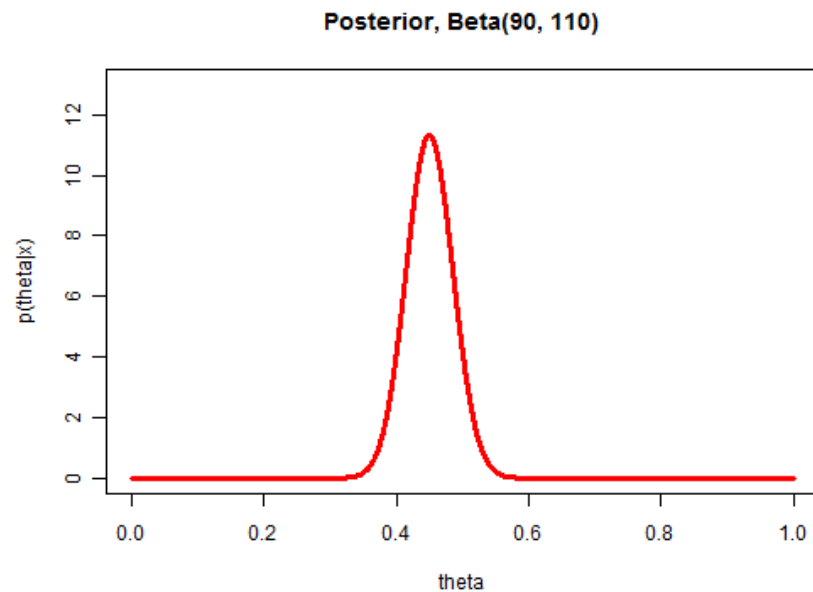

$$\hat{\theta}_{\text{MLE}} = \frac{40}{100}$$

$$\hat{\theta}_{\text{MAP}} = \frac{42}{104}$$

$$\hat{\theta}_{\text{Bayes}} = \frac{43}{106}$$

$$n = 100, n_1 = 40$$

Strong prior ($\alpha=50, \beta=50$)



$$\hat{\theta}_{\text{MLE}} = \frac{40}{100}$$

$$\hat{\theta}_{\text{MAP}} = \frac{89}{198}$$

$$\hat{\theta}_{\text{Bayes}} = \frac{90}{200}$$

$$n = 100, n_1 = 40$$

Based on a sufficient statistic

Let X_1, \dots, X_n be a sample from a Bernoulli (θ) population, $0 < \theta < 1$.

Let $Y = \sum_{i=1}^n X_i$ be the sum of 1s (a sufficient statistic of θ). Then,

$$p(\theta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1} (1 - \theta)^{\beta-1}, \quad p(y | \theta) = \binom{n}{y} \theta^y (1 - \theta)^{n-y},$$

$$\begin{aligned} p(y) &= \int_0^1 p(\theta) p(y | \theta) d\theta = \int_0^1 \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1} (1 - \theta)^{\beta-1} \binom{n}{y} \theta^y (1 - \theta)^{n-y} d\theta \\ &= \binom{n}{y} \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \frac{\Gamma(\alpha + y)\Gamma(\beta + n - y)}{\Gamma(\alpha + \beta + n)}. \end{aligned}$$

Therefore,

$$p(\theta | y) = \frac{p(\theta) p(y | \theta)}{p(y)} = \frac{\Gamma(\alpha + \beta + n)}{\Gamma(\alpha + y)\Gamma(\beta + n - y)} \theta^{(\alpha+y)-1} (1 - \theta)^{(\beta+n-y)-1}.$$

Hence,

$$\hat{\theta} = E(\theta | y) = \frac{\alpha + y}{\alpha + \beta + n} = \underbrace{\left(\frac{\alpha + \beta}{\alpha + \beta + n} \right)}_{\text{Prior weight}} \underbrace{\left(\frac{\alpha}{\alpha + \beta} \right)}_{\text{Prior mean}} + \underbrace{\left(\frac{n}{\alpha + \beta + n} \right)}_{\text{Sample weight}} \underbrace{\left(\frac{y}{n} \right)}_{\text{Sample mean}}$$

Multinomial likelihood

Let X_1, \dots, X_n be iid random variables from a multinomial trial ($\boldsymbol{\theta}$), where $\boldsymbol{\theta} = (\theta_1, \dots, \theta_k)$, $0 < \theta_k < 1$ and $\sum_{k=1}^m \theta_k = 1$. Then the count vector $\mathbf{n} = (n_1, \dots, n_k)$, as a sufficient statistic, has a multinomial distribution with cell probability $\boldsymbol{\theta}$.

Therefore

$$p(\mathbf{n} \mid \boldsymbol{\theta}) = \frac{(\sum_{k=1}^m n_k)!}{\prod_{k=1}^m n_k!} \prod_{k=1}^m \theta_k^{n_k}$$

Dirichlet prior

The likelihood is

$$p(\mathbf{n} \mid \theta) \propto \prod_{k=1}^m \theta_k^{n_k}$$

The conjugate prior is the Dirichlet distribution

$$p(\theta) = \frac{\Gamma(\sum_{k=1}^m \alpha_k)}{\prod_{k=1}^m \Gamma(\alpha_k)} \prod_{k=1}^m \theta_k^{\alpha_k - 1}$$

The parameters in the prior are called **hyper-parameters**

Integrate out the parameter

$$\begin{aligned} p(\mathbf{x}) &= \int p(\boldsymbol{\theta}) p(\mathbf{x} \mid \boldsymbol{\theta}) d\boldsymbol{\theta} \\ &= \int \left[\frac{\Gamma(\sum_{k=1}^m \alpha_k)}{\prod_{k=1}^m \Gamma(\alpha_k)} \prod_{k=1}^m \theta_k^{\alpha_k - 1} \right] \left[\frac{(\sum_{k=1}^m n_k)!}{\prod_{k=1}^m n_k!} \prod_{k=1}^m \theta_k^{n_k} \right] d\boldsymbol{\theta} \\ &= \frac{(\sum_{k=1}^m n_k)!}{\prod_{k=1}^m n_k!} \frac{\Gamma(\sum_{k=1}^m \alpha_k)}{\prod_{k=1}^m \Gamma(\alpha_k)} \underbrace{\int \prod_{k=1}^m \theta_k^{\alpha_k + n_k - 1} d\boldsymbol{\theta}}_{= \left[\frac{\Gamma(\sum_{k=1}^m (\alpha_k + n_k))}{\prod_{k=1}^m \Gamma(\alpha_k + n_k)} \right]^{-1}} \\ &= \frac{(\sum_{k=1}^m n_k)!}{\prod_{k=1}^m n_k!} \frac{\Gamma(\sum_{k=1}^m \alpha_k)}{\prod_{k=1}^m \Gamma(\alpha_k)} \frac{\prod_{k=1}^m \Gamma(\alpha_k + n_k)}{\Gamma(\sum_{k=1}^m (\alpha_k + n_k))} \end{aligned}$$

Obtain the posterior

$$\begin{aligned} p(\theta \mid \mathbf{x}) &= \frac{p(\theta)p(\mathbf{x} \mid \theta)}{p(\mathbf{x})} \\ &= \frac{\left[\frac{\Gamma(\sum_{k=1}^m \alpha_k)}{\prod_{k=1}^m \Gamma(\alpha_k)} \prod_{k=1}^m \theta^{\alpha_k - 1} \right] \left[\frac{(\sum_{k=1}^m n_k)!}{\prod_{k=1}^m n_k!} \prod_{k=1}^m \theta^{n_k} \right]}{\frac{(\sum_{k=1}^m n_k)!}{\prod_{k=1}^m n_k!} \frac{\Gamma(\sum_{k=1}^m \alpha_k)}{\prod_{k=1}^m \Gamma(\alpha_k)} \frac{\prod_{k=1}^m \Gamma(\alpha_k + n_k)}{\Gamma(\sum_{k=1}^m (\alpha_k + n_k))}} \\ &= \frac{\Gamma(\sum_{k=1}^m (\alpha_k + n_k))}{\prod_{k=1}^m \Gamma(\alpha_k + n_k)} \prod_{k=1}^m \theta^{\alpha_k + n_k} \end{aligned}$$

Multinomial Bayes estimate

Therefore

$$\begin{aligned}\hat{\theta}_k &= E(\theta_k \mid \mathbf{n}) \\&= \frac{\alpha_k + n_k}{\sum_{k=1}^m (\alpha_k + n_k)} \\&= \underbrace{\frac{\sum_{k=1}^m \alpha_k}{\sum_{k=1}^m \alpha_k + \sum_{k=1}^m n_k}}_{\text{Prior weight}} \underbrace{\frac{\alpha_k}{\sum_{k=1}^m \alpha_k}}_{\text{Prior mean}} + \underbrace{\frac{\sum_{k=1}^m n_k}{\sum_{k=1}^m \alpha_k + \sum_{k=1}^m n_k}}_{\text{Sample weight}} \underbrace{\frac{n_k}{\sum_{k=1}^m n_k}}_{\text{Sample mean}}\end{aligned}$$

Posterior mean is the weighted average of the prior mean and the sample mean

Normal likelihood (variance known)

Let X_1, \dots, X_n be a sample from a $N(\mu, \sigma^2)$ population, where σ^2 is known. we have the likelihood function as

$$p(\mathbf{x} \mid \mu) = (2\pi\sigma^2)^{-\frac{n}{2}} \exp\left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2\right]$$
$$\propto \exp\left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2\right]$$

Normal prior

The likelihood is

$$p(\mathbf{x} \mid \mu) \propto \exp \left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The conjugate prior is also the normal distribution

$$\mu \sim N(\xi, \tau^2)$$

$$p(\mu \mid \xi, \tau^2) \propto \exp \left[-\frac{1}{2\tau^2} (\mu - \xi)^2 \right]$$

The parameters in the prior are called **hyper-parameters**

$$p(\mu \mid \mathbf{x}) \propto p(\mu, \mathbf{x}) = p(\mu)p(\mathbf{x} \mid \mu)$$

Prior \times Likelihood

$$p(\mu \mid \xi, \tau) p(\mathbf{x} \mid \mu)$$

$$\propto \exp\left[-\frac{1}{2\tau^2}(\mu - \xi)^2\right] \times \exp\left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2\right]$$

$$\propto \exp\left[-\frac{1}{2\tau^2}(\mu^2 - 2\xi\mu + \xi^2) - \frac{1}{2\sigma^2}\left(n\mu^2 - 2\left(\sum_{i=1}^n x_i\right)\mu + \sum_{i=1}^n x_i^2\right)\right]$$

$$\propto \exp\left[-\frac{1}{2}\left(\left(\frac{1}{\tau^2} + \frac{n}{\sigma^2}\right)\mu^2 - 2\left(\frac{\xi}{\tau^2} + \frac{n\bar{x}}{\sigma^2}\right)\mu + \left(\frac{\xi^2}{\tau^2} + \frac{\sum_{i=1}^n x_i^2}{\sigma^2}\right)\right)\right]$$

$$\propto \exp\left[-\frac{1}{2}\left(\frac{1}{\tau^2} + \frac{n}{\sigma^2}\right)\left(\mu^2 - 2\frac{\sigma^2\xi + \tau^2 n\bar{x}}{\sigma^2 + \tau^2 n}\mu + \left(\frac{\sigma^2\xi + \tau^2 n\bar{x}}{\sigma^2 + \tau^2 n}\right)^2\right)\right]$$

$$\propto \exp\left[-\frac{1}{2}\left(\frac{\sigma^2\tau^2}{\sigma^2 + \tau^2 n}\right)^{-1}\left(\mu - \frac{\sigma^2\xi + \tau^2 n\bar{x}}{\sigma^2 + \tau^2 n}\right)^2\right]$$

Obtain the posterior

$$p(\mu \mid \mathbf{x}) \propto \exp \left[-\frac{1}{2} \left(\frac{\sigma^2 \tau^2}{\sigma^2 + \tau^2 n} \right)^{-1} \left(\mu - \frac{\sigma^2 \xi + \tau^2 n \bar{x}}{\sigma^2 + \tau^2 n} \right)^2 \right]$$


$$\mu \mid \mathbf{x} \sim N \left(\frac{\sigma^2 \xi + \tau^2 n \bar{x}}{\sigma^2 + \tau^2 n}, \frac{\sigma^2 \tau^2}{\sigma^2 + \tau^2 n} \right)$$

Posterior is also a normal distribution


Prior (variance=1, known)




With one observation




With two observations




With three observations




With five observations




With ten observations




With 20 observations




With 50 observations



With 100 observations



With 1000 observations



Any problem?

Not a full conjugate prior, because the variance appears in the posterior but not in the prior

$$\mu \sim N(\xi, \tau^2)$$

$$\mu \mid \mathbf{x} \sim N\left(\frac{\sigma^2 \xi + \tau^2 n \bar{x}}{\sigma^2 + \tau^2 n}, \frac{\sigma^2 \tau^2}{\sigma^2 + \tau^2 n}\right)$$

Normal likelihood (precision known)

Let X_1, \dots, X_n be a sample from a $N(\mu, \lambda^{-1})$ population, where $\lambda = 1 / \sigma^2$ is known, we have the likelihood function as

$$p(\mathbf{x} \mid \mu) = (2\pi)^{-\frac{n}{2}} \lambda^{\frac{n}{2}} \exp \left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$
$$\propto \exp \left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

$\lambda = 1 / \sigma^2$ is called the **precision**

Normal prior

The likelihood is

$$p(\mathbf{x} \mid \mu) = \sqrt{\frac{\lambda}{2\pi}} \exp \left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The conjugate prior is also the normal distribution

$$\mu \mid \lambda \sim N(\xi, (\kappa\lambda)^{-1})$$

$$p(\mu \mid \xi, \kappa\lambda) \propto \exp \left[-\frac{\kappa\lambda}{2} (\mu - \xi)^2 \right]$$

The parameters in the prior are called **hyper-parameters**

$$p(\mu \mid \mathbf{x}) \propto p(\mu, \mathbf{x}) = p(\mu)p(\mathbf{x} \mid \mu)$$

Prior \times Likelihood

$$p(\mu \mid \xi, \kappa\lambda)p(\mathbf{x} \mid \mu)$$

$$\propto \exp\left[-\frac{\kappa\lambda}{2}(\mu - \xi)^2\right] \times \exp\left[-\frac{\lambda}{2}\sum_{i=1}^n (x_i - \mu)^2\right]$$

$$\propto \exp\left[-\frac{\kappa\lambda}{2}(\mu^2 - 2\xi\mu + \xi^2) - \frac{\lambda}{2}\left(n\mu^2 - 2\left(\sum_{i=1}^n x_i\right)\mu + \sum_{i=1}^n x_i^2\right)\right]$$

$$\propto \exp\left[-\frac{\lambda}{2}\left((\kappa + n)\mu^2 - 2(\kappa\xi + n\bar{x})\mu + (\kappa\xi^2 + \sum_{i=1}^n x_i^2)\right)\right]$$

$$\propto \exp\left[-\frac{(\kappa + n)\lambda}{2}\left(\mu^2 - 2\frac{\kappa\xi + n\bar{x}}{\kappa + n}\mu + \left(\frac{\kappa\xi + n\bar{x}}{\kappa + n}\right)^2\right)\right]$$

$$\propto \exp\left[-\frac{(\kappa + n)\lambda}{2}\left(\mu - \frac{\kappa\xi + n\bar{x}}{\kappa + n}\right)^2\right]$$

Obtain the posterior

$$p(\mu \mid \mathbf{x}) \propto \exp \left[-\frac{(\kappa + n)\lambda}{2} \left(\mu - \frac{\kappa\xi + n\bar{x}}{\kappa + n} \right)^2 \right]$$

$$\mu \mid \mathbf{x} \sim N(\tilde{\xi}, (\tilde{\kappa}\lambda)^{-1})$$

$$\tilde{\xi} = \frac{\kappa\xi + n\bar{x}}{\kappa + n}$$

$$\tilde{\kappa} = \kappa + n$$

Posterior is also a normal distribution

Normal Bayes estimate

Therefore

$$\begin{aligned}\hat{\mu} &= E(\mu \mid \mathbf{x}) \\ &= \frac{\kappa \xi + n \bar{x}}{\kappa + n} \\ &= \underbrace{\frac{\kappa}{\kappa + n}}_{\text{Prior weight}} \underbrace{\xi}_{\text{Prior mean}} + \underbrace{\frac{n}{\kappa + n}}_{\text{Sample weight}} \underbrace{\bar{x}}_{\text{Sample mean}}\end{aligned}$$

Posterior mean is the weighted average of the prior mean and the sample mean

How about the precision?

Let X_1, \dots, X_n be a sample from a $N(\mu, \lambda^{-1})$ population, where μ is known, we have the likelihood function as

$$\begin{aligned} p(\mathbf{x} \mid \lambda) &= (2\pi)^{-\frac{n}{2}} \lambda^{\frac{n}{2}} \exp\left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2\right] \\ &\propto \lambda^{\frac{n}{2}} \exp\left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2\right] \end{aligned}$$

$\lambda = 1 / \sigma^2$ is called the **precision**

Gamma distribution

Pdf

$$f(x \mid \text{shape}=\alpha, \text{scale}=\theta) = \frac{1}{\Gamma(\alpha)\theta^\alpha} x^{\alpha-1} e^{-x/\theta}, 0 \leq x < \infty, \alpha > 0, \theta > 0$$

$$EX = \alpha\theta, \quad \text{Var} X = \alpha\theta^2$$

Pdf

$$f(x \mid \text{shape}=\alpha, \text{rate}=\beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}, 0 \leq x < \infty, \alpha > 0, \beta > 0$$

$$EX = \frac{\alpha}{\beta}, \quad \text{Var} X = \frac{\alpha}{\beta^2}$$

Gamma prior

The likelihood is

$$p(\mathbf{x} \mid \lambda) \propto \lambda^{\frac{n}{2}} \exp \left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The natural prior is the Gamma distribution

$$\lambda \sim \text{Gamma}(\alpha, \text{rate} = \beta)$$

$$p(\lambda \mid \alpha, \text{rate} = \beta) \propto \lambda^{\alpha-1} \exp(-\beta\lambda)$$

Note that in the definition of the Gamma distribution, we use rate to take the place of the scale. The parameters in the prior are called hyper-parameters.

$$p(\lambda, \mathbf{x}) = p(\lambda)p(\mathbf{x} \mid \lambda)$$

Prior \times Likelihood

$$p(\lambda \mid \alpha, \beta) p(\mathbf{x} \mid \lambda)$$

$$\propto \lambda^{\alpha-1} \exp(-\beta\lambda) \times \lambda^{\frac{n}{2}} \exp\left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2\right]$$

$$\propto \lambda^{\alpha + \frac{n}{2} - 1} \exp\left[-\lambda \left(\beta + \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2\right)\right]$$

$$\propto \lambda^{\tilde{\alpha}-1} \exp(-\lambda\tilde{\beta})$$

$$\tilde{\alpha} = \alpha + \frac{n}{2}$$

$$\tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2$$

Obtain the posterior

$$p(\lambda \mid \mathbf{x}) \propto \lambda^{\tilde{\alpha}-1} \exp(-\lambda\tilde{\beta})$$

$$\lambda \mid \mathbf{x} \sim \text{Gamma}(\tilde{\alpha}, \text{rate} = \tilde{\beta})$$

$$\tilde{\alpha} = \alpha + \frac{n}{2}$$

$$\tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2$$

Posterior is also a Gamma distribution

How about the variance?

Let X_1, \dots, X_n be a sample from a $N(\mu, \sigma^2)$ population, where μ is known, we have the likelihood function as

$$p(\mathbf{x} \mid \sigma^2) = (2\pi)^{-\frac{n}{2}} (\sigma^2)^{-\frac{n}{2}} \exp \left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

$$\propto \left(\frac{1}{\sigma^2} \right)^{\frac{n}{2}} \exp \left[-\frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

Inverse-gamma distribution

Gamma pdf

$$f(x \mid \text{shape}=\alpha, \text{rate}=\beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}, 0 \leq x < \infty, \alpha > 0, \beta > 0$$

$$EX = \alpha / \beta, \text{Var} X = \alpha / \beta^2$$

The transformation $1/X$ will yield an inverse-gamma distribution with pdf

$$f(x \mid \alpha, \beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} (1/x)^{\alpha+1} e^{-\beta/x}, 0 \leq x < \infty, \alpha > 0, \beta > 0$$

$$EX = \frac{\beta}{\alpha - 1}, \quad \text{Var} X = \frac{\beta^2}{(\alpha - 1)^2(\alpha - 2)}$$

Inverse gamma prior

The likelihood is

$$p(\mathbf{x} \mid \sigma^2) \propto \left(\frac{1}{\sigma^2} \right)^{\frac{n}{2}} \exp \left[-\frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The conjugate prior is the Inverse gamma distribution

$$p(\sigma^2 \mid \alpha, \beta) \propto (1 / \sigma^2)^{\alpha+1} \exp(-\beta / \sigma^2)$$

The parameters in the prior are called **hyper-parameters**.

$$p(\sigma^2, \mathbf{x}) = p(\sigma^2) p(\mathbf{x} \mid \sigma^2)$$

Prior \times Likelihood

$$p(\sigma^2 \mid \alpha, \beta) p(\mathbf{x} \mid \sigma^2)$$

$$\propto (1 / \sigma^2)^{\alpha+1} \exp(-\beta / \sigma^2) \times (1 / \sigma^2)^{\frac{n}{2}} \exp\left[-\frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2\right]$$

$$\propto (1 / \sigma^2)^{\alpha + \frac{n}{2} + 1} \exp\left[-\frac{1}{\sigma^2} \left(\beta + \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2\right)\right]$$

$$\propto (1 / \sigma^2)^{\tilde{\alpha}+1} \exp(-\tilde{\beta} / \sigma^2)$$

$$\tilde{\alpha} = \alpha + \frac{n}{2}$$

$$\tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2$$

Obtain the posterior

$$p(\sigma^2 \mid \mathbf{x}) \propto (1 / \sigma^2)^{\tilde{\alpha}-1} \exp(-\tilde{\beta} / \sigma^2)$$

$$\sigma^2 \mid \mathbf{x} \sim \text{Inverse} - \text{Gamma}(\tilde{\alpha}, \tilde{\beta})$$

$$\tilde{\alpha} = \alpha + \frac{n}{2}$$

$$\tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2$$

Posterior is also a Inverse-Gamma distribution

Inverse-chi-square distribution

Gamma pdf

$$f_{\text{gamma}}(x \mid \text{shape}=\alpha, \text{rate}=\beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}, 0 \leq x < \infty, \alpha > 0, \beta > 0$$

The special case ($\alpha=p/2, \beta=1/2$) is called chi-squared distribution

$$f_{\chi_p^2}(x \mid p) = \frac{1}{\Gamma(p/2)2^{p/2}} x^{p/2-1} e^{-x/2}, 0 \leq x < \infty, p > 0$$

The transformation $1/X$ on the chi-squared distribution will yield an inverse-chi-squared distribution. Or equivalently, the special case ($\alpha=p/2, \beta=1/2$) of an inverse-gamma distribution will yield an inverse-chi-squared distribution.

$$f_{\text{inv-}\chi_p^2}(x \mid p) = \frac{1}{\Gamma(p/2)2^{p/2}} (1/x)^{p/2+1} e^{-1/(2x)}, 0 \leq x < \infty, p > 0$$

Scaled-inverse-chi-square distribution

Inverse-gamma pdf

$$f_{inv-gamma}(x \mid \alpha, \beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} (1/x)^{\alpha+1} e^{-\beta/x}, 0 \leq x < \infty, \alpha > 0, \beta > 0$$

The special case ($\alpha=p/2, \beta=pq/2$) of an inverse-gamma distribution will yield a scaled-inverse-chi-square distribution. Alternatively, the transformation pq/X on the chi-squared distribution will also yield a scaled-inverse-chi-squared distribution.

$$f_{scaled-inv-\chi^2}(x \mid p, q) = \frac{(pq/2)^{p/2}}{\Gamma(p/2)} (1/x)^{p/2+1} e^{-(pq)/(2x)},$$

$$0 \leq x < \infty, p > 0, q > 0$$

Scaled-inverse-chi-squared prior

The likelihood is

$$p(\mathbf{x} \mid \sigma^2) \propto \left(\frac{1}{\sigma^2} \right)^{\frac{n}{2}} \exp \left[-\frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The natural prior is the Scaled-inverse-chi-square distribution

$$p(\sigma^2 \mid \eta, \tau^2) \propto (1 / \sigma^2)^{\eta/2+1} \exp[-\eta\tau^2 / (2\sigma^2)]$$

The parameters in the prior are called **hyper-parameters**.

$$p(\sigma^2, \mathbf{x}) = p(\sigma^2)p(\mathbf{x} \mid \sigma^2)$$

Prior \times Likelihood

$$p(\sigma^2 \mid \eta, \tau^2) p(\mathbf{x} \mid \sigma^2)$$

$$\propto (1 / \sigma^2)^{\eta/2+1} \exp[-\eta\tau^2 / (2\sigma^2)] \times (1 / \sigma^2)^{\frac{n}{2}} \exp\left[-\frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2\right]$$

$$\propto (1 / \sigma^2)^{\frac{\eta+n}{2}+1} \exp\left[-\frac{1}{2\sigma^2} \left(\eta\tau^2 + n \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2\right)\right]$$

$$\propto (1 / \sigma^2)^{\frac{\eta+n}{2}+1} \exp\left[-\frac{1}{2\sigma^2} (\eta\tau^2 + ns_n^2)\right]$$

$$\propto (1 / \sigma^2)^{\tilde{\eta}/2+1} \exp[-\tilde{\eta}\tilde{\tau}^2 / (2\sigma^2)]$$

$$\tilde{\eta} = \eta + n$$

$$\tilde{\tau}^2 = \frac{\eta\tau^2 + ns_n^2}{\eta + n}$$

Obtain the posterior

$$p(\sigma^2 \mid \mathbf{x}) \propto (1 / \sigma^2)^{\tilde{\eta}+1} \exp(-\tilde{\eta}\tilde{\tau}^2 / (2\sigma^2))$$

$$\sigma \mid \mathbf{x} \sim \textit{Scaled-inverse-chi-square}(\tilde{\eta}, \tilde{\tau}^2)$$

$$\tilde{\eta} = \eta + n$$

$$\tilde{\tau}^2 = \frac{\eta\tau^2 + ns_n^2}{\eta + n}$$

Posterior is also a Scaled-inverse-chi-square distribution

For both parameters

Let X_1, \dots, X_n be a sample from a $N(\mu, \lambda^{-1})$ population, we have the likelihood function as

$$\begin{aligned} p(\mathbf{x} \mid \mu, \lambda) &= (2\pi)^{-\frac{n}{2}} \lambda^{\frac{n}{2}} \exp\left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2\right] \\ &\propto \lambda^{\frac{n}{2}} \exp\left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2\right] \end{aligned}$$

$\lambda = 1 / \sigma^2$ is called the **precision**

Normal-Gamma prior

The likelihood is

$$p(\mathbf{x} \mid \mu, \lambda) \propto \lambda^{\frac{n}{2}} \exp \left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The conjugate prior for the mean is the Normal distribution

$$\mu \mid \lambda \sim N(\xi, (\kappa\lambda)^{-1}) \Rightarrow p(\mu \mid \xi, \kappa\lambda) \propto \lambda^{\frac{1}{2}} \exp \left[-\frac{\kappa\lambda}{2} (\mu - \xi)^2 \right]$$

The conjugate prior for the precision is the Gamma distribution

$$\lambda \sim \text{Gamma}(\alpha, \beta) \Rightarrow p(\lambda \mid \alpha, \beta) \propto \lambda^{\alpha-1} \exp(-\beta\lambda)$$

The parameters in the prior are called **hyper-parameters**.

$$p(\mu, \lambda \mid \mathbf{x}) \propto p(\mu, \lambda, \mathbf{x}) = p(\lambda)p(\mu \mid \lambda)p(\mathbf{x} \mid \mu, \lambda)$$

Normal-Gamma distribution

The product

$$\begin{aligned} p(\mu, \lambda \mid \xi, \kappa, \alpha, \beta) &= p(\mu \mid \lambda, \xi, \kappa, \alpha, \beta) p(\lambda \mid \xi, \kappa, \alpha, \beta) \\ &\propto \lambda^{\frac{1}{2}} \exp\left[-\frac{\kappa\lambda}{2}(\mu - \xi)^2\right] \lambda^{\alpha-1} \exp(-\beta\lambda) \end{aligned}$$

is the kernel of the so called **Normal-Gamma** distribution.

In order to determine the constant, consider the integral

$$\int_0^\infty \int_{-\infty}^\infty \lambda^{\frac{1}{2}} \exp\left[-\frac{\kappa\lambda}{2}(\mu - \xi)^2\right] \lambda^{\alpha-1} \exp(-\beta\lambda) d\mu d\lambda$$
$$\left(\frac{\kappa}{2\pi}\right)^{-\frac{1}{2}} \left(\frac{\beta^\alpha}{\Gamma(\alpha)}\right)^{-1}$$

Determine the constant

$$\begin{aligned} & \int_0^\infty \int_{-\infty}^\infty \lambda^{\frac{1}{2}} \exp\left[-\frac{\kappa\lambda}{2}(\mu - \xi)^2\right] \lambda^{\alpha-1} \exp(-\beta\lambda) d\mu d\lambda \\ &= \int_0^\infty \left\{ \int_{-\infty}^\infty \lambda^{\frac{1}{2}} \exp\left[-\frac{\kappa\lambda}{2}(\mu - \xi)^2\right] d\mu \right\} \lambda^{\alpha-1} \exp(-\beta\lambda) d\lambda \\ &= \int_0^\infty \left\{ \left(\frac{\kappa}{2\pi}\right)^{-\frac{1}{2}} \int_{-\infty}^\infty \left(\frac{\kappa\lambda}{2\pi}\right)^{\frac{1}{2}} \exp\left[-\frac{\kappa\lambda}{2}(\mu - \xi)^2\right] d\mu \right\} \lambda^{\alpha-1} \exp(-\beta\lambda) d\lambda \\ &= \left(\frac{\kappa}{2\pi}\right)^{-\frac{1}{2}} \left(\frac{\beta^\alpha}{\Gamma(\alpha)}\right)^{-1} \int_0^\infty \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} \exp(-\beta\lambda) d\lambda \\ &= \left(\frac{\kappa}{2\pi}\right)^{-\frac{1}{2}} \left(\frac{\beta^\alpha}{\Gamma(\alpha)}\right)^{-1} \end{aligned}$$

Therefore, the constant is $\frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi}\right)^{\frac{1}{2}}$, and the Normal-Gamma pdf is

$$\frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi}\right)^{\frac{1}{2}} \lambda^{\frac{1}{2}} \exp\left[-\frac{\kappa\lambda}{2}(\mu - \xi)^2\right] \lambda^{\alpha-1} \exp(-\beta\lambda)$$

Marginal distribution

$$\begin{aligned} p(\lambda) &= \int_{-\infty}^{\infty} \frac{\beta^{\alpha}}{\Gamma(\alpha)} \left(\frac{\kappa\lambda}{2\pi} \right)^{\frac{1}{2}} \exp \left[-\frac{\kappa\lambda}{2} (\mu - \xi)^2 \right] \lambda^{\alpha-1} \exp(-\beta\lambda) d\mu \\ &= \frac{\beta^{\alpha}}{\Gamma(\alpha)} \lambda^{\alpha-1} \exp(-\beta\lambda) \left\{ \int_{-\infty}^{\infty} \left(\frac{\kappa\lambda}{2\pi} \right)^{\frac{1}{2}} \exp \left[-\frac{\kappa\lambda}{2} (\mu - \xi)^2 \right] d\mu \right\} \\ &= \frac{\beta^{\alpha}}{\Gamma(\alpha)} \lambda^{\alpha-1} \exp(-\beta\lambda) \end{aligned}$$

$$p(\mu \mid \lambda) = \left(\frac{\kappa\lambda}{2\pi} \right)^{\frac{1}{2}} \exp \left[-\frac{\kappa\lambda}{2} (\mu - \xi)^2 \right]$$

$$\lambda \sim \text{Gamma}(\alpha, \text{rate} = \beta)$$

$$\mu \mid \lambda \sim \text{Normal}(\xi, (\kappa\lambda)^{-1})$$

Marginal distribution

$$\begin{aligned}
 p(\mu) &= \int_0^\infty \frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi} \right)^{\frac{1}{2}} \lambda^{\frac{1}{2}} \exp \left[-\frac{\kappa\lambda}{2} (\mu - \xi)^2 \right] \lambda^{\alpha-1} \exp(-\beta\lambda) d\lambda \\
 &= \frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi} \right)^{\frac{1}{2}} \int_0^\infty \lambda^{(\alpha+1/2)-1} \exp \left[-\lambda \left(\beta + \frac{\kappa}{2} (\mu - \xi)^2 \right) \right] d\lambda \\
 &= \frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi} \right)^{\frac{1}{2}} \left[\frac{(\beta + (\mu - \xi)^2 \kappa / 2)^{\alpha+1/2}}{\Gamma(\alpha + 1/2)} \right]^{-1} \\
 &= \frac{\Gamma\left(\frac{2\alpha+1}{2}\right)}{\Gamma\left(\frac{2\alpha}{2}\right)} \left(\frac{\kappa}{2\beta\pi} \right)^{\frac{1}{2}} \left(1 + \frac{\kappa(\mu - \xi)^2}{2\beta} \right)^{-(2\alpha+1)/2} \\
 &= \frac{\Gamma\left(\frac{2\alpha+1}{2}\right)}{\Gamma\left(\frac{2\alpha}{2}\right)} \frac{1}{\sqrt{(\beta / \alpha\kappa)(2\alpha)\pi}} \left(1 + \frac{1}{2\alpha} \left(\frac{\mu - \xi}{\sqrt{\beta / (\alpha\kappa)}} \right)^2 \right)^{-(2\alpha+1)/2}
 \end{aligned}$$

Let $t = \frac{\mu - \xi}{\sqrt{\beta / (\alpha\kappa)}}$, $\mu = \sqrt{\beta / (\alpha\kappa)}t + \xi$, $\frac{d\mu}{dt} = \sqrt{\beta / (\alpha\kappa)}$

$$p(t) = \frac{\Gamma\left(\frac{2\alpha+1}{2}\right)}{\Gamma\left(\frac{2\alpha}{2}\right)} \frac{1}{\sqrt{(2\alpha)\pi}} \left(1 + \frac{t^2}{2\alpha} \right)^{-(2\alpha+1)/2}, \text{ which is a } t_{2\alpha} \text{ pdf.}$$

Marginal distribution

In summary,

If

$$\mu, \lambda \sim \text{Normal-Gamma}(\xi, \kappa, \alpha, \beta)$$

Then

$$\lambda \sim \text{Gamma}(\alpha, \beta)$$

$$\mu \sim \text{Shifted (and scaled) } Student's\ t \text{ with } 2\alpha \text{ degrees of freedom}$$

$$\text{Location parameter} = \xi$$

$$\text{Scale parameter} = \sqrt{\beta / (\alpha \kappa)}$$

$$\mu \mid \lambda \sim N(\xi, (\kappa \lambda)^{-1})$$

Prior \times Likelihood

$$p(\mu \mid \xi, \kappa \lambda) p(\lambda \mid \alpha, \beta) p(\mathbf{x} \mid \mu, \lambda)$$

$$\propto \lambda^{\frac{1}{2}} \exp\left[-\frac{\kappa \lambda}{2} (\mu - \xi)^2\right] \lambda^{\alpha-1} \exp(-\beta \lambda) \lambda^{\frac{n}{2}} \exp\left[-\frac{\lambda}{2} \sum_{i=1}^n (x_i - \mu)^2\right]$$

$$\propto \lambda^{\frac{1}{2}} \lambda^{\alpha+\frac{n}{2}-1} \exp\left[-\frac{\lambda}{2} \left(2\beta + \kappa(\mu - \xi)^2 + \sum_{i=1}^n (x_i - \mu)^2\right)\right]$$

$$\propto \lambda^{\frac{1}{2}} \lambda^{\alpha+\frac{n}{2}-1} \exp\left[-\frac{\lambda}{2} \left(2\beta + (\kappa + n)\mu^2 - 2(\kappa\xi + n\bar{x})\mu + (\kappa\xi^2 + \sum_{i=1}^n x_i^2)\right)\right]$$

$$\propto \lambda^{\frac{1}{2}} \lambda^{\alpha+\frac{n}{2}-1} \exp\left[-\frac{(\kappa + n)\lambda}{2} \left(\mu - \frac{\kappa\xi + n\bar{x}}{\kappa + n}\right)^2 - \lambda \left(\beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa(\bar{x} - \xi)^2}{\kappa + n}\right)\right]$$

$$\propto \lambda^{\frac{1}{2}} \exp\left[-\frac{\tilde{\kappa} \lambda}{2} (\mu - \tilde{\xi})^2\right] \lambda^{\tilde{\alpha}-1} \exp(-\lambda \tilde{\beta})$$

$$\tilde{\xi} = \frac{\kappa\xi + n\bar{x}}{\kappa + n}, \tilde{\kappa} = \kappa + n, \tilde{\alpha} = \alpha + \frac{n}{2}, \tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2$$

Obtain the posterior

$$p(\mu, \lambda \mid \mathbf{x}) \propto \lambda^{\frac{1}{2}} \exp\left[-\frac{\tilde{\kappa}\lambda}{2}(\mu - \tilde{\xi})^2\right] \lambda^{\tilde{\alpha}-1} \exp(-\lambda\tilde{\beta})$$

$$\mu, \lambda \mid \mathbf{x} \sim \text{Normal-Gamma}(\tilde{\xi}, \tilde{\kappa}, \tilde{\alpha}, \tilde{\beta})$$

$$\tilde{\xi} = \frac{\kappa\xi + n\bar{x}}{\kappa + n}, \quad \tilde{\kappa} = \kappa + n, \quad \tilde{\alpha} = \alpha + \frac{n}{2},$$

$$\tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2$$

Posterior is also a Normal-Gamma distribution

Bayes estimate

$$\hat{\mu}_{\text{Bayes}} = \tilde{\xi} = \frac{\kappa\xi + n\bar{x}}{\kappa + n}$$

$$\hat{\lambda}_{\text{Bayes}} = \frac{\tilde{\alpha}}{\tilde{\beta}} = \frac{2\alpha + n}{2\beta + \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2}$$

$$\hat{\sigma}_{\text{Bayes}}^2 = \frac{1}{\hat{\lambda}_{\text{Bayes}}}$$

Normal-Inverse-gamma prior

The likelihood is

$$p(\mathbf{x} \mid \mu, \sigma^2) \propto (1 / \sigma^2)^{\frac{n}{2}} \exp \left[-\frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The conjugate prior for the mean is the Normal distribution

$$\mu \mid \sigma^2 \sim N(\xi, \sigma^2 / \kappa) \Rightarrow p(\mu \mid \xi, \sigma^2 / \kappa) \propto (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2} \right]$$

The conjugate prior for the variance is the Inverse-gamma distribution

$$\sigma^2 \sim \text{Inverse - gamma}(\alpha, \beta) \Rightarrow p(\sigma^2 \mid \alpha, \beta) \propto (1 / \sigma^2)^{\alpha+1} \exp(-\beta / \sigma^2)$$

The product is a **Normal-Inverse-gamma distribution**

$$\mu, \sigma^2 \sim N(\xi, \sigma^2 / \kappa) \text{Inverse - gamma}(\alpha, \beta)$$

Normal-Inverse-gamma distribution

The product

$$\begin{aligned} p(\mu, \lambda \mid \xi, \kappa, \alpha, \beta) &= p(\mu \mid \sigma^2, \xi, \kappa, \alpha, \beta) p(\sigma^2 \mid \xi, \kappa, \alpha, \beta) \\ &\propto (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2} \right] (1 / \sigma^2)^{\alpha-1} \exp(-\beta / \sigma^2) \end{aligned}$$

is the kernel of the so called **Normal-Inverse-gamma** distribution. In order to determine the constant, consider the integral

$$\int_0^\infty \int_{-\infty}^\infty (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2} \right] (1 / \sigma^2)^{\alpha+1} \exp(-\beta / \sigma^2) d\mu d\sigma^2$$

Determine the constant

$$\begin{aligned}
 & \int_0^\infty \int_{-\infty}^\infty (1/\sigma^2)^{\frac{1}{2}} \exp\left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2}\right] (1/\sigma^2)^{\alpha+1} \exp(-\beta/\sigma^2) d\mu d\sigma^2 \\
 &= \int_0^\infty \left\{ \int_{-\infty}^\infty (1/\sigma^2)^{\frac{1}{2}} \exp\left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2}\right] d\mu \right\} (1/\sigma^2)^{\alpha+1} \exp(-\beta/\sigma^2) d\lambda \\
 &= \int_0^\infty \left\{ \left(\frac{\kappa}{2\pi}\right)^{-\frac{1}{2}} \int_{-\infty}^\infty \left(\frac{1}{2\pi(\sigma^2/\kappa)}\right)^{\frac{1}{2}} \exp\left[-\frac{(\mu - \xi)^2}{2(\sigma^2/\kappa)}\right] d\mu \right\} \lambda^{\alpha-1} \exp(-\beta\lambda) d\lambda \\
 &= \left(\frac{\kappa}{2\pi}\right)^{-\frac{1}{2}} \left(\frac{\beta^\alpha}{\Gamma(\alpha)}\right)^{-1} \int_0^\infty \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} \exp(-\beta\lambda) d\lambda \\
 &= \left(\frac{\kappa}{2\pi}\right)^{-\frac{1}{2}} \left(\frac{\beta^\alpha}{\Gamma(\alpha)}\right)^{-1}
 \end{aligned}$$

Therefore, the constant is $\frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi}\right)^{\frac{1}{2}}$, and the Normal-Inverse-gamma pdf is

$$\frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi}\right)^{\frac{1}{2}} (1/\sigma^2)^{\frac{1}{2}} \exp\left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2}\right] (1/\sigma^2)^{\alpha+1} \exp(-\beta/\sigma^2)$$

Prior \times Likelihood

$$p(\mu \mid \xi, \sigma^2 / \kappa) p(\sigma^2 \mid \alpha, \beta) p(\mathbf{x} \mid \mu, \sigma^2)$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2} \right] (1 / \sigma^2)^{\alpha+1} \exp(-\beta / \sigma^2) (1 / \sigma^2)^{\frac{n}{2}} \exp \left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} (1 / \sigma^2)^{\alpha+\frac{n}{2}+1} \exp \left[-\frac{1}{2\sigma^2} \left(2\beta + \kappa(\mu - \xi)^2 + \sum_{i=1}^n (x_i - \mu)^2 \right) \right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} (1 / \sigma^2)^{\alpha+\frac{n}{2}+1} \exp \left[-\frac{1}{2\sigma^2} \left(2\beta + (\kappa + n)\mu^2 - 2(\kappa\xi + n\bar{x})\mu + (\kappa\xi^2 + \sum_{i=1}^n x_i^2) \right) \right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} (1 / \sigma^2)^{\alpha+\frac{n}{2}+1} \exp \left[-\frac{(\kappa + n)}{2\sigma^2} \left(\mu - \frac{\kappa\xi + n\bar{x}}{\kappa + n} \right)^2 - \frac{1}{\sigma^2} \left(\beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa(\bar{x} - \xi)^2}{\kappa + n} \right) \right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\tilde{\kappa}}{2\sigma^2} (\mu - \tilde{\xi})^2 \right] (1 / \sigma^2)^{\tilde{\alpha}+1} \exp(-\tilde{\beta} / \sigma^2)$$

$$\tilde{\xi} = \frac{\kappa\xi + n\bar{x}}{\kappa + n}, \tilde{\kappa} = \kappa + n, \tilde{\alpha} = \alpha + \frac{n}{2}, \tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2$$

Obtain the posterior

$$p(\mu, \sigma^2 \mid \mathbf{x}) \propto (1 / \sigma^2)^{\frac{1}{2}} \exp\left[-\frac{\tilde{\kappa}}{2\sigma^2}(\mu - \tilde{\xi})^2\right] (1 / \sigma^2)^{\tilde{\alpha}+1} \exp(-\tilde{\beta} / \sigma^2)$$

$$\mu, \sigma^2 \mid \mathbf{x} \sim \text{Normal - Inverse - gamma}(\tilde{\xi}, \tilde{\kappa}, \tilde{\alpha}, \tilde{\beta})$$

$$\tilde{\xi} = \frac{\kappa \xi + n \bar{x}}{\kappa + n}, \quad \tilde{\kappa} = \kappa + n, \quad \tilde{\alpha} = \alpha + \frac{n}{2},$$

$$\tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2$$

Posterior is also a Normal-Inverse-Gamma distribution

Normal-Scaled-inverse-chi-square prior

The likelihood is

$$p(\mathbf{x} \mid \mu, \sigma^2) \propto (1 / \sigma^2)^{\frac{n}{2}} \exp \left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right]$$

The conjugate prior for the mean is the Normal distribution

$$\mu \mid \sigma^2 \sim \text{N}(\xi, \sigma^2 / \kappa) \Rightarrow p(\mu \mid \xi, \sigma^2 / \kappa) \propto (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\kappa}{2\sigma^2} (\mu - \xi)^2 \right]$$

The conjugate prior for the variance is the Scaled-inverse-chi-square distribution

$$\sigma^2 \sim \text{Scaled-Inverse-chi-square}(\eta, \tau^2) \Rightarrow p(\sigma^2 \mid \eta, \tau^2) \propto (1 / \sigma^2)^{\eta/2+1} \exp \left[-\frac{\eta\tau^2}{2\sigma^2} \right]$$

The product is a **Normal-Scaled-inverse-chi-square distribution**

$$\mu, \sigma^2 \sim \text{N}(\xi, \sigma^2 / \kappa) \text{Scaled-inverse-chi-square}(\eta, \tau^2)$$

Normal-Scaled-inverse-chi-square distribution

The product

$$\begin{aligned} p(\mu, \sigma^2 \mid \xi, \kappa, \eta, \tau^2) &= p(\mu \mid \sigma^2, \xi, \kappa, \eta, \tau^2) p(\sigma^2 \mid \xi, \kappa, \eta, \tau^2) \\ &\propto (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2} \right] (1 / \sigma^2)^{\eta/2+1} \exp \left[-\frac{\eta \tau^2}{2\sigma^2} \right] \end{aligned}$$

is the kernel of the so called **Normal-Scaled-inverse-chi-square** distribution. In order to determine the constant, consider the integral

$$\int_0^\infty \int_{-\infty}^\infty (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\kappa}{\sigma^2} \frac{(\mu - \xi)^2}{2} \right] (1 / \sigma^2)^{\eta/2+1} \exp \left[-\frac{\eta \tau^2}{2\sigma^2} \right] d\mu d\sigma^2$$

Prior \times Likelihood

$$p(\mu \mid \xi, \sigma^2 / \kappa) p(\sigma^2 \mid \eta, \tau^2) p(\mathbf{x} \mid \mu, \sigma^2)$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} \exp\left[-\frac{\kappa}{2\sigma^2}(\mu - \xi)^2\right] (1 / \sigma^2)^{\eta/2+1} \exp\left[-\frac{\eta\tau^2}{2\sigma^2}\right] (1 / \sigma^2)^{\frac{n}{2}} \exp\left[-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2\right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} (1 / \sigma^2)^{\frac{\eta+n}{2}+1} \exp\left[-\frac{1}{2\sigma^2} \left(\eta\tau^2 + \kappa(\mu - \xi)^2 + \sum_{i=1}^n (x_i - \mu)^2 \right)\right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} (1 / \sigma^2)^{\frac{\eta+n}{2}+1} \exp\left[-\frac{1}{2\sigma^2} \left(\eta\tau^2 + (\kappa + n)\mu^2 - 2(\kappa\xi + n\bar{x})\mu + (\kappa\xi^2 + \sum_{i=1}^n x_i^2) \right)\right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} (1 / \sigma^2)^{\frac{\eta+n}{2}+1} \exp\left[-\frac{(\kappa + n)}{2\sigma^2} \left(\mu - \frac{\kappa\xi + n\bar{x}}{\kappa + n} \right)^2 - \frac{1}{2\sigma^2} \left(\eta\tau^2 + \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{n\kappa(\bar{x} - \xi)^2}{\kappa + n} \right)\right]$$

$$\propto (1 / \sigma^2)^{\frac{1}{2}} \exp\left[-\frac{\tilde{\kappa}}{2\sigma^2}(\mu - \tilde{\xi})^2\right] (1 / \sigma^2)^{\tilde{\alpha}+1} \exp\left[-\frac{\tilde{\eta}\tilde{\tau}^2}{2\sigma^2}\right]$$

$$\tilde{\xi} = \frac{\kappa\xi + n\bar{x}}{\kappa + n}, \tilde{\kappa} = \kappa + n, \tilde{\eta} = \eta + n, \tilde{\eta}\tilde{\tau}^2 = \eta\tau^2 + \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{n\kappa}{\kappa + n}(\bar{x} - \xi)^2$$

Obtain the posterior

$$p(\mu, \sigma^2 \mid \mathbf{x}) \propto (1 / \sigma^2)^{\frac{1}{2}} \exp \left[-\frac{\tilde{\kappa}}{\sigma^2} \frac{(\mu - \tilde{\xi})^2}{2} \right] (1 / \sigma^2)^{\tilde{\alpha}+1} \exp(-\tilde{\eta} \tilde{\tau}^2 / \sigma^2)$$

$$\mu, \sigma^2 \mid \mathbf{x} \sim \text{Normal-Scaled-inverse-chi-square}(\tilde{\xi}, \tilde{\kappa}, \tilde{\eta}, \tilde{\tau}^2)$$

$$\tilde{\xi} = \frac{\kappa \xi + n \bar{x}}{\kappa + n}, \quad \tilde{\kappa} = \kappa + n, \quad \tilde{\eta} = \eta + n,$$

$$\tilde{\eta} \tilde{\tau}^2 = \eta \tau^2 + \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{n \kappa}{\kappa + n} (\bar{x} - \xi)^2$$

Posterior is also a Normal-Scaled-inverse-chi-square distribution

Making predictions

- Using the MLE estimate as the value of the parameter and then calculate the density for a given value of the observation
- Using the MAP estimate as the value of the parameter to calculate the density
- Using the Bayes estimate as the value of the parameter to calculate the density
- All of these approaches are not pure Bayesian ways
- A pure Bayesian way will average over all possible parameters

$$p(x^* \mid \mathbf{x}) = \int p(x^* \mid \theta) p(\theta \mid \mathbf{x}) d\theta$$

Bernoulli sample

Posterior distribution

$$\theta \mid \mathbf{x} \sim \text{Beta}(\alpha + n_1, \beta + n - n_1)$$

Likelihood of a single observation

$$p(x \mid \theta) = \theta^x (1 - \theta)^{1-x}$$

Integral

$$\begin{aligned} \int p(x \mid \theta) p(\theta \mid \mathbf{x}) d\theta &= \int \theta^x (1 - \theta)^{1-x} \frac{\Gamma(\alpha + \beta + n)}{\Gamma(\alpha + n_1) \Gamma(\beta + n - n_1)} \theta^{\alpha + n_1 - 1} (1 - \theta)^{\beta + n - n_1 - 1} d\theta \\ &= \frac{\Gamma(\alpha + \beta + n)}{\Gamma(\alpha + n_1) \Gamma(\beta + n - n_1)} \frac{\Gamma(\alpha + n_1 + x) \Gamma(\beta + n - n_1 + 1 - x)}{\Gamma(\alpha + \beta + n + 1)} \\ &= \frac{\Gamma(\alpha + \beta + n)}{\Gamma(\alpha + \beta + n + 1)} \frac{\Gamma(\alpha + n_1 + x)}{\Gamma(\alpha + n_1)} \frac{\Gamma(\beta + n - n_1 + 1 - x)}{\Gamma(\beta + n - n_1)} \\ &= \left[\frac{\alpha + n_1}{\alpha + \beta + n} \right]^x \left[\frac{\beta + n - n_1}{\alpha + \beta + n} \right]^{1-x} \end{aligned}$$

Marginal likelihood

- In MLE, the value of the likelihood function measures how good the model is fitted
- In Bayesian, the marginal likelihood provide a similar measure
- Marginal likelihood can also be used to make predictions

$$p(\mathbf{x}) = \int p(\theta)p(\mathbf{x} \mid \theta)d\theta$$

Bernoulli marginal likelihood

Marginal likelihood

$$\begin{aligned} p(\mathbf{x}) &= \int p(\theta) p(\mathbf{x} \mid \theta) d\theta \\ &= \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \frac{\Gamma(\alpha + n_1)\Gamma(\beta + n - n_1)}{\Gamma(\alpha + \beta + n)} \\ &= \left[\frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \right] \left[\frac{\Gamma(\tilde{\alpha} + \tilde{\beta})}{\Gamma(\tilde{\alpha})\Gamma(\tilde{\beta})} \right]^{-1} \end{aligned}$$

$$\tilde{\alpha} = \alpha + n_1$$

$$\tilde{\beta} = \beta + n - n_1$$

Make predictions

Now, we notice that

$$\begin{aligned} p(x \mid \mathbf{x}) &= \int p(\theta \mid \mathbf{x}) p(x \mid \theta) d\theta \\ &= \left[\frac{\Gamma(\tilde{\alpha} + \tilde{\beta})}{\Gamma(\tilde{\alpha})\Gamma(\tilde{\beta})} \right] \left[\frac{\Gamma(\hat{\alpha} + \hat{\beta})}{\Gamma(\hat{\alpha})\Gamma(\hat{\beta})} \right]^{-1} \end{aligned}$$

$$\tilde{\alpha} = \alpha + n_1$$

$$\tilde{\beta} = \beta + n - n_1$$

$$\hat{\alpha} = \tilde{\alpha} + x$$

$$\hat{\beta} = \tilde{\beta} + 1 - x$$

Normal marginal likelihood

Marginal likelihood

$$\begin{aligned}
 p(\mathbf{x}) &= \int p(\lambda) p(\mu \mid \lambda) p(\mathbf{x} \mid \mu, \lambda) d\mu d\lambda \\
 &= \underbrace{\left[\frac{\beta^\alpha}{\Gamma(\alpha)} \left(\frac{\kappa}{2\pi} \right)^{\frac{1}{2}} \right]}_{\text{Prior constant}} \underbrace{\left[\left(\frac{1}{2\pi} \right)^{\frac{n}{2}} \right]}_{\text{Likelihood constant}} \underbrace{\left[\frac{\tilde{\beta}^{\tilde{\alpha}}}{\Gamma(\tilde{\alpha})} \left(\frac{\tilde{\kappa}}{2\pi} \right)^{\frac{1}{2}} \right]^{-1}}_{\text{Inverse of the posterior constant}}
 \end{aligned}$$

$$\tilde{\xi} = \frac{\kappa \xi + n \bar{x}}{\kappa + n}, \tilde{\kappa} = \kappa + n, \tilde{\alpha} = \alpha + \frac{n}{2}, \tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2$$

Make a single prediction

Marginal likelihood

$$\begin{aligned}
 p(x \mid \mathbf{x}) &= \int p(\mu, \lambda \mid \mathbf{x}) p(x \mid \mu, \lambda) d\mu d\lambda \\
 &= \underbrace{\left[\frac{\tilde{\beta}^{\tilde{\alpha}}}{\Gamma(\tilde{\alpha})} \left(\frac{\tilde{\kappa}}{2\pi} \right)^{\frac{1}{2}} \right]}_{\text{Posterior constant}} \underbrace{\left[\left(\frac{1}{2\pi} \right)^{\frac{1}{2}} \right]}_{\text{Likelihood constant}} \underbrace{\left[\frac{\hat{\beta}^{\hat{\alpha}}}{\Gamma(\hat{\alpha})} \left(\frac{\hat{\kappa}}{2\pi} \right)^{\frac{1}{2}} \right]^{-1}}_{\text{Inverse of the updated posterior constant}}
 \end{aligned}$$

$$\begin{aligned}
 \tilde{\xi} &= \frac{\kappa \xi + n \bar{x}}{\kappa + n}, \quad \tilde{\kappa} = \kappa + n, \quad \tilde{\alpha} = \alpha + \frac{n}{2}, \quad \tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2 \\
 \hat{\xi} &= \frac{\tilde{\kappa} \tilde{\xi} + x}{\tilde{\kappa} + 1}, \quad \hat{\kappa} = \tilde{\kappa} + 1, \quad \hat{\alpha} = \tilde{\alpha} + \frac{1}{2}, \quad \beta = \tilde{\beta} + \frac{1}{2} \frac{\tilde{\kappa}}{\tilde{\kappa} + 1} (x - \tilde{\xi})^2
 \end{aligned}$$

Make multiple predictions

Marginal likelihood

$$\begin{aligned}
 p(\mathbf{x}^* \mid \mathbf{x}) &= \int p(\mu, \lambda \mid \mathbf{x}) p(\mathbf{x}^* \mid \mu, \lambda) d\mu d\lambda \\
 &= \underbrace{\left[\frac{\tilde{\beta}^{\tilde{\alpha}}}{\Gamma(\tilde{\alpha})} \left(\frac{\tilde{\kappa}}{2\pi} \right)^{\frac{1}{2}} \right]}_{\text{Posterior constant}} \underbrace{\left[\left(\frac{1}{2\pi} \right)^{\frac{m}{2}} \right]}_{\text{Likelihood constant}} \underbrace{\left[\frac{\hat{\beta}^{\hat{\alpha}}}{\Gamma(\hat{\alpha})} \left(\frac{\hat{\kappa}}{2\pi} \right)^{\frac{1}{2}} \right]^{-1}}_{\text{Inverse of the updated posterior constant}}
 \end{aligned}$$

$$\begin{aligned}
 \tilde{\xi} &= \frac{\kappa \xi + n \bar{x}}{\kappa + n}, \quad \tilde{\kappa} = \kappa + n, \quad \tilde{\alpha} = \alpha + \frac{n}{2}, \quad \tilde{\beta} = \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{1}{2} \frac{n\kappa}{\kappa + n} (\bar{x} - \xi)^2 \\
 \hat{\xi} &= \frac{\tilde{\kappa} \tilde{\xi} + m \bar{x}^*}{\tilde{\kappa} + m}, \quad \hat{\kappa} = \tilde{\kappa} + m, \quad \hat{\alpha} = \tilde{\alpha} + \frac{m}{2}, \quad \hat{\beta} = \tilde{\beta} + \frac{1}{2} \sum_{i=1}^n (x_i^* - \bar{x}^*)^2 + \frac{1}{2} \frac{m\tilde{\kappa}}{\tilde{\kappa} + m} (\bar{x}^* - \tilde{\xi})^2
 \end{aligned}$$